

Multitasking and Concurrency - 1

Review of theoretical background

Prof. Himanshu B. Dave
e-Infochips, Ahmedabad

June, 2010 / e-Infochips, Ahmedabad

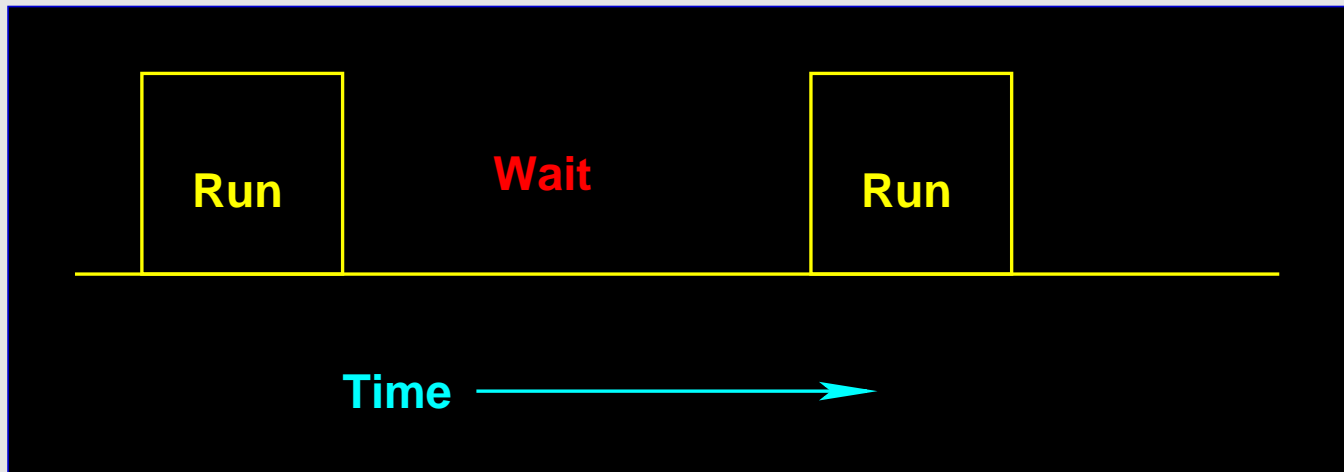


Contents

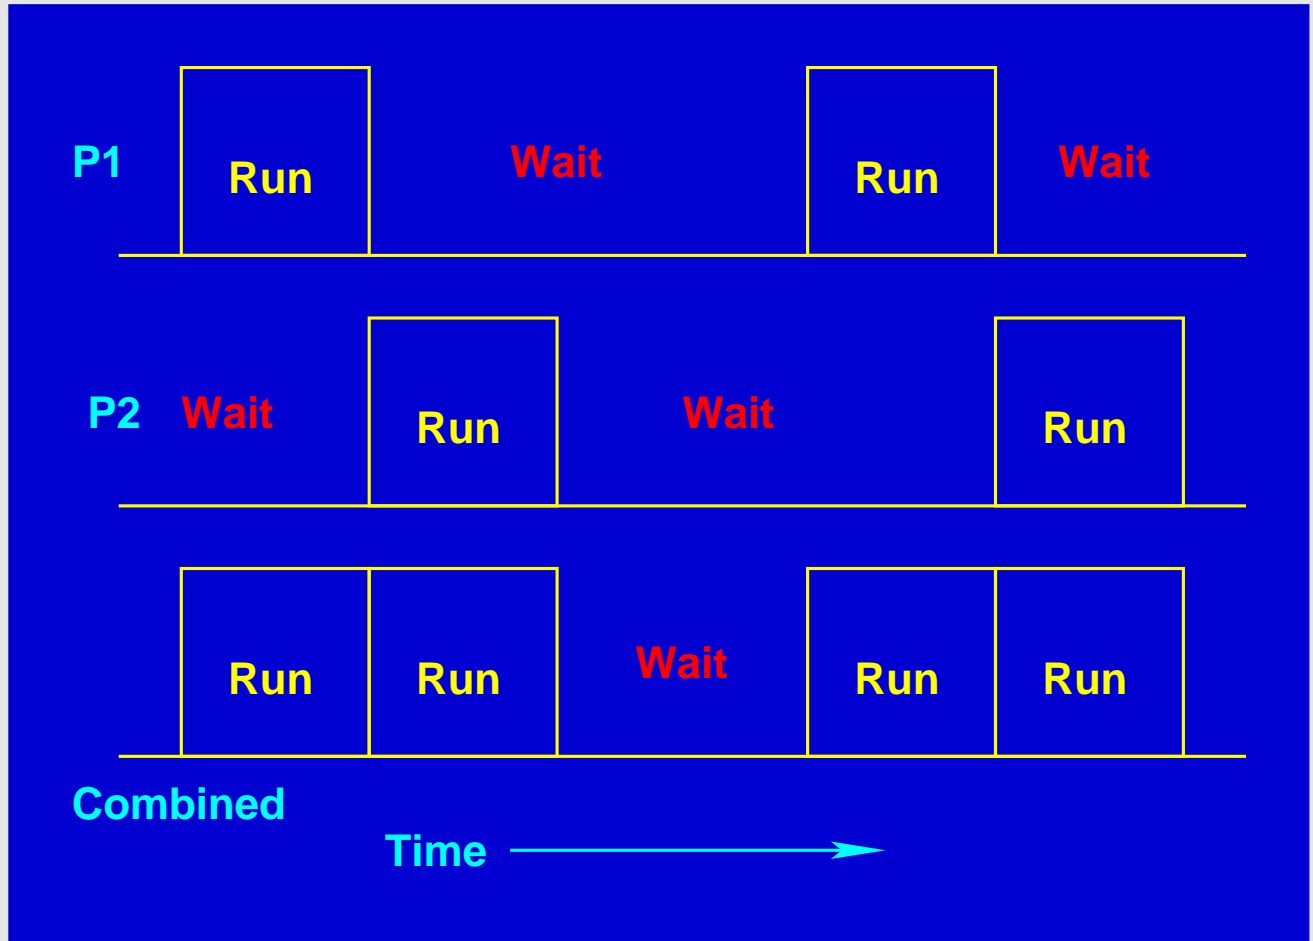
- ▶ What is Concurrency?
- ▶ Concurrency vs. Parallelism
- ▶ Need and Advantages
- ▶ Difficulties with concurrent execution
- ▶ Concurrency by processes and threads
- ▶ Sharing of Resources
- ▶ Synchronization Techniques
- ▶ Threads - Life Cycle
- ▶ On writing concurrent code




Uniprogramming



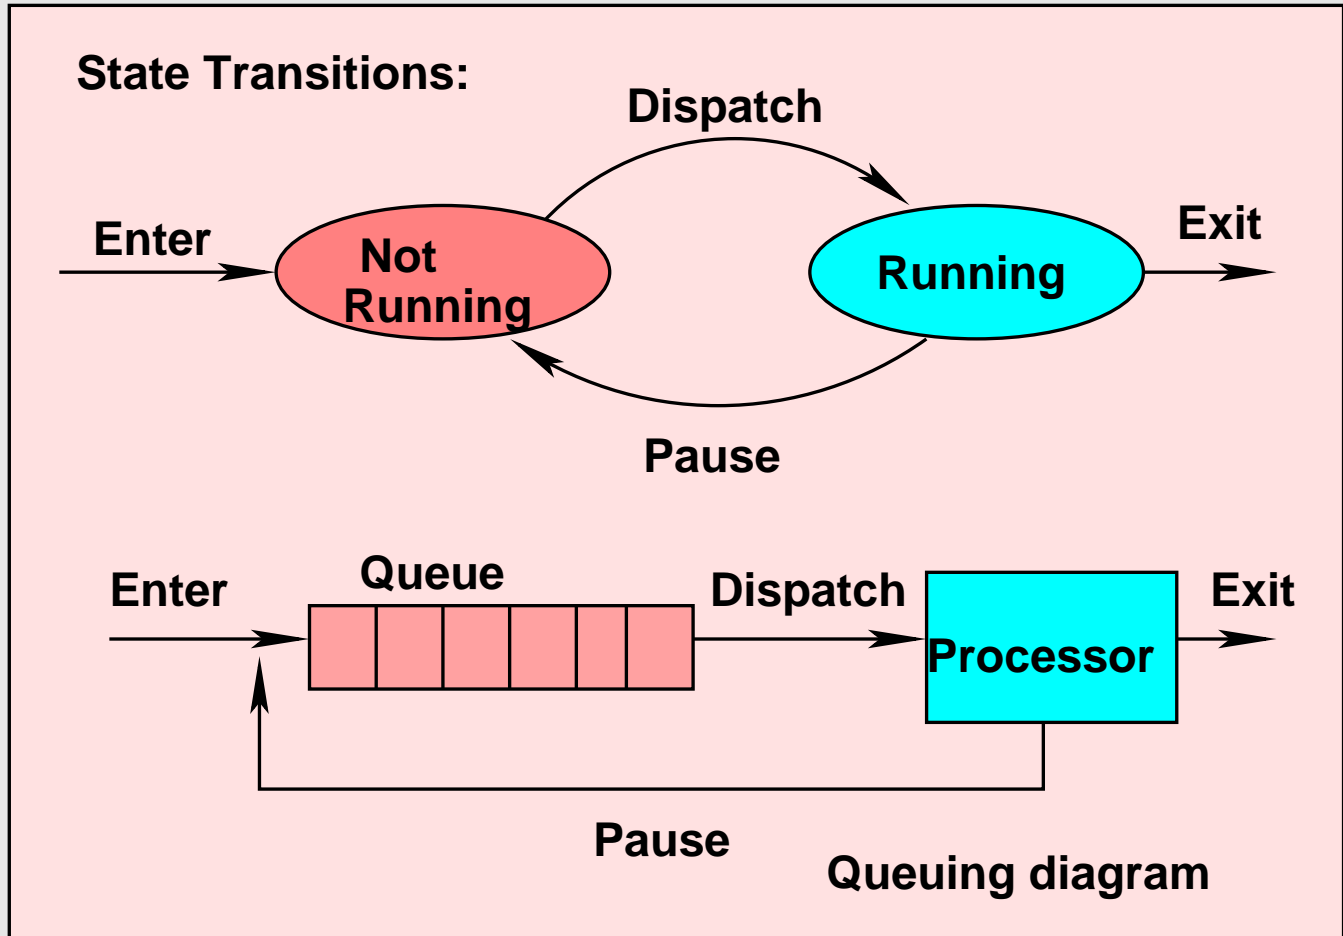
Multi-programming



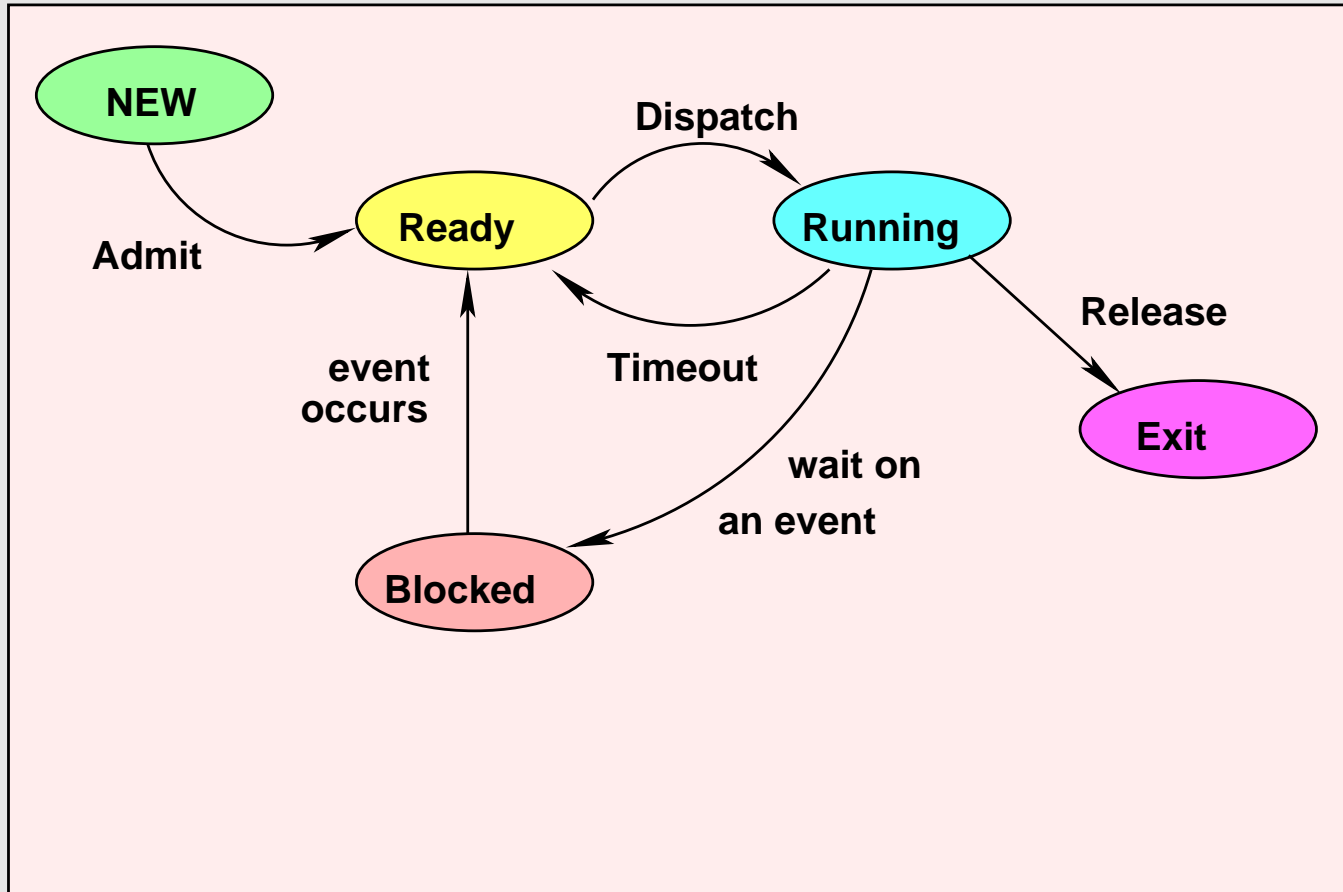
Several Tasks executing on a Single Processor



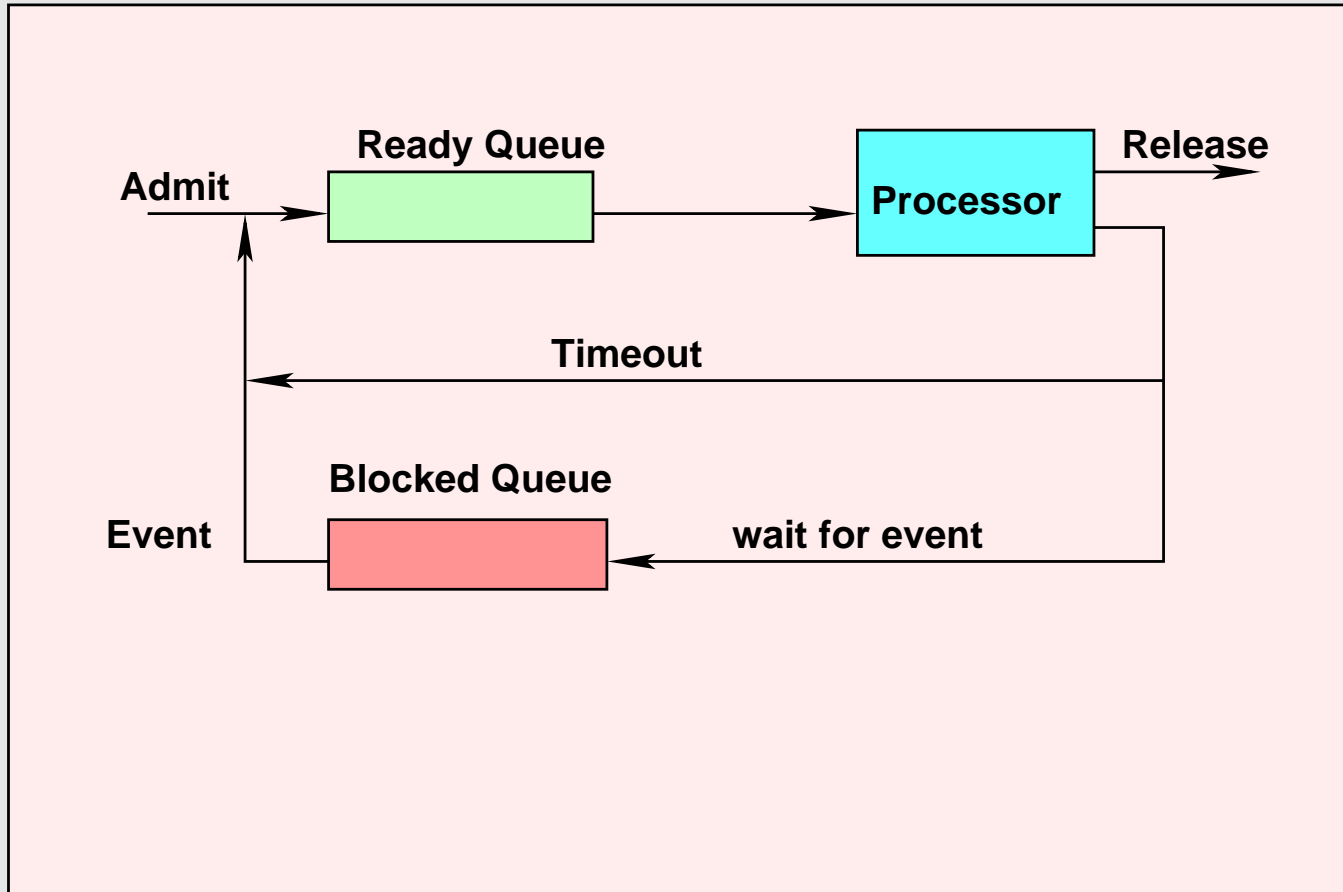
Two-state model of a Process



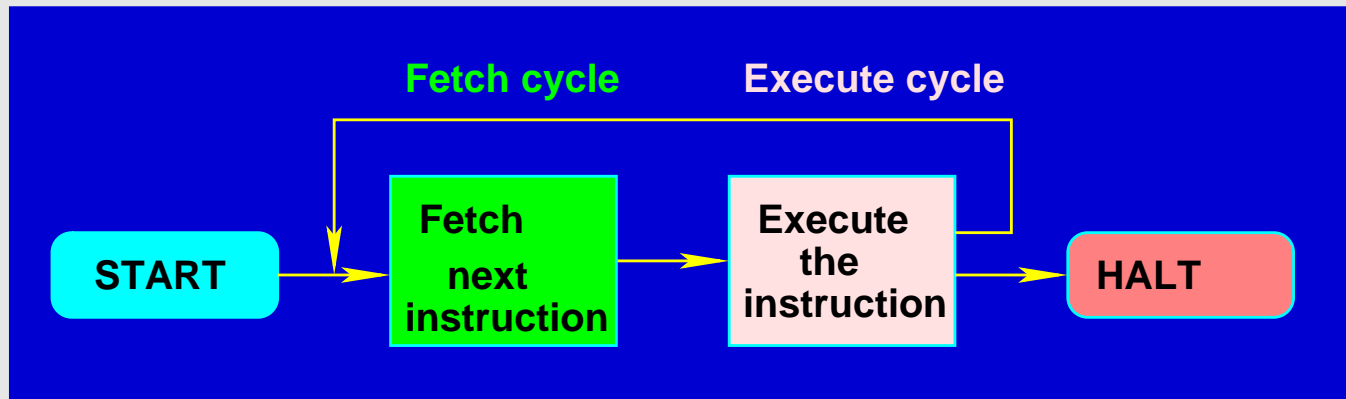
Five-state model of a Process



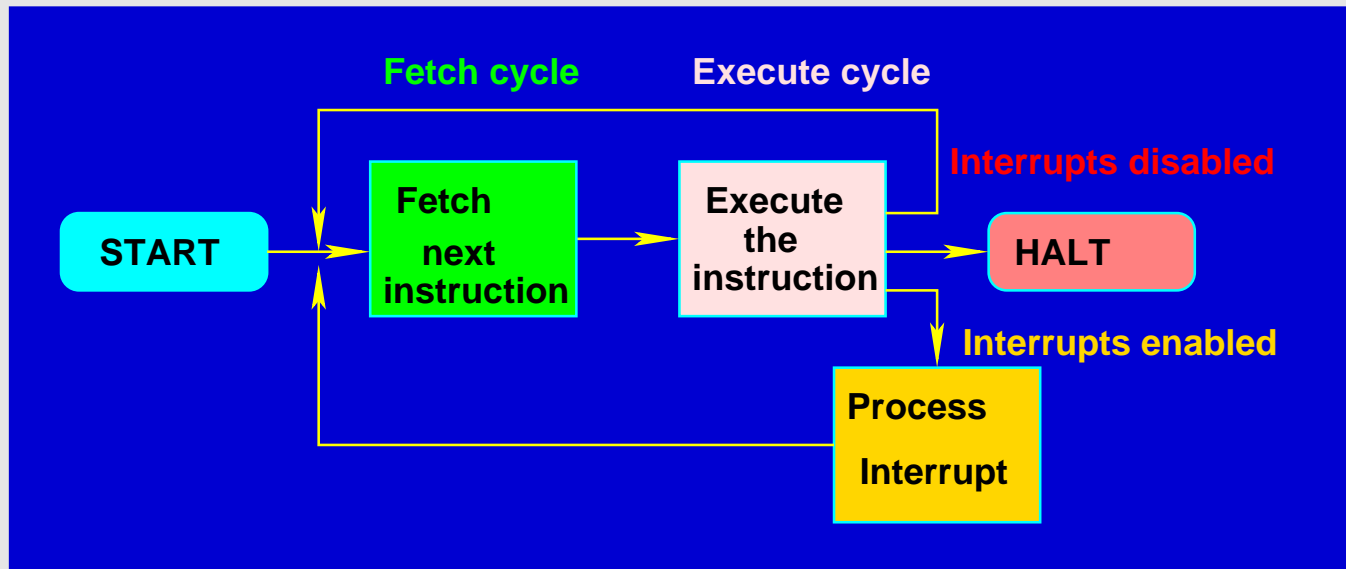
Queuing model for five states of a Process



How atomicity is created: F/E cycle



How atomicity is created: Interrupt cycle



Concurrent Processing

What is Concurrency?

- ▶ "sharing of resources in the same time frame".
- ▶ several processes share the same CPU, memory or I/O devices;
- ▶ requires correct co-ordination of processes;
- ▶ previously O/S handled concurrency,
- ▶ no more true, now programmers have also to worry about concurrency:
 - complex applications,
 - multi-processor architecture,
 - distributed systems;

programming to achieve and control concurrency



Concurrency

- ▶ A **sequential** program has a **single** *thread of control* or execution-context. Its execution is called a **process**.
- ▶ A **concurrent** program has **multiple** *threads of control*, or execution-contexts. These *may be* executed as **parallel processes**.



Concurrency

What is "going on at the same time"?

processes progress to their completion at the same time

Is it parallelism?

No: parallelism would mean instant by instant, simultaneous operations by more than one processors

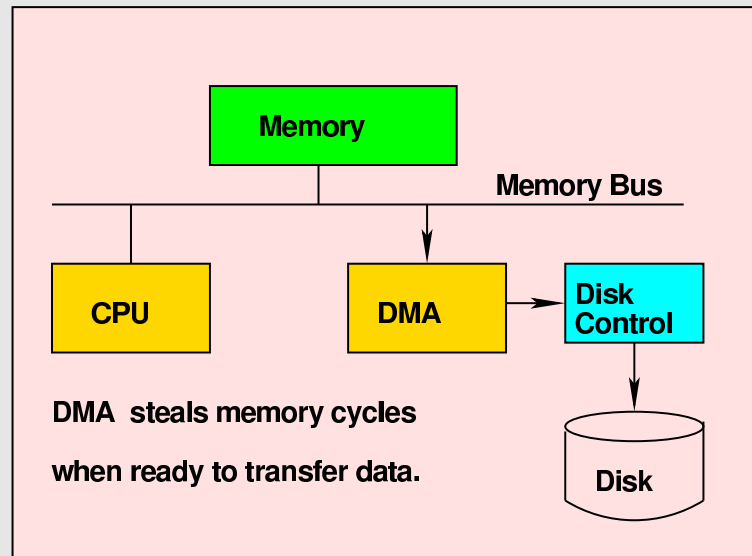


Figure 1

Concurrency and Parallelism

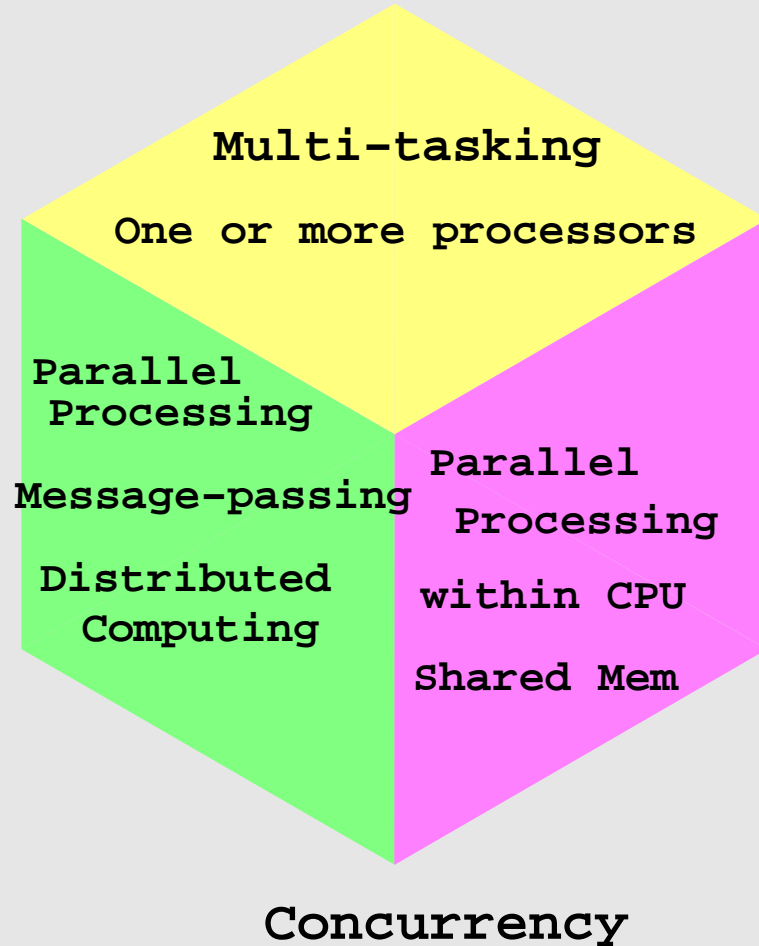
A concurrent program can be executed by:

- ▶ **Multiprogramming:** processes *share* *one* or more processors
- ▶ **Multiprocessing:** each process runs on its own processor but with shared memory
- ▶ **Distributed processing:** each process runs on its own processor connected by a communication network to others

Assume only that all processes make positive finite progress.



Concurrency can be found in:



Multi-processing: bus with shared memory





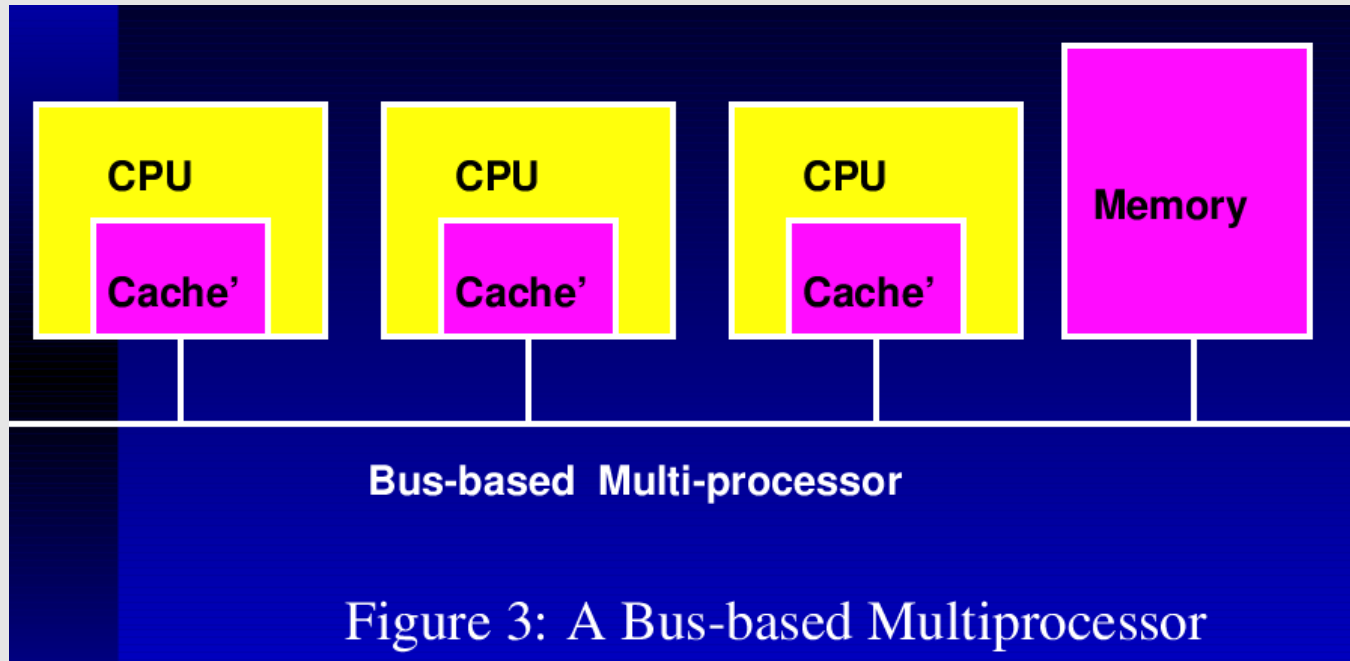
Figure 1: Bus-based, Shared Memory




Multi-processing: bus with private memory



Multi-processing: independent CPU's on bus



Multi-processing: PU and MEM connected by Cross-bar switch



Where and when is Concurrency needed?

- ▶ logic of application demands it
- ▶ application is easier to visualize and implement that way;
- ▶ resource utilization issues demand it;
- ▶ speed-up execution by parallel operation on multiple processors, **this has become a necessity as we reach limits of VLSI component size**
- ▶ coordinate distributed services (e.g. Web-Services)




Amdahl's Law

Amdahl's Law

If 50% of your application is parallel and 50% is serial, you can't get more than a factor of 2 speedup, no matter how many processors it runs on.*

*In general, if a fraction α of an application can be run in parallel and the rest must run serially, the speedup is at most $1/(1-\alpha)$.



Gene M. Amdahl

But, whose application can be decomposed into just a serial part and a parallel part? For *my* application, what speedup should I expect?




Two Pioneer Computer Engineers

Gordon E. Moore

Gene Amdahl

No, they're not actors



Moore's and Amdahl's Laws

- **Moore's Law.**

- *The number of transistors that can be inexpensively placed on an integrated circuit is increasing exponentially.*
- Not true anymore!

- **Amdahl's Law.**

- *Performance decreases as number of processors increases once there is even a small percentage of non-parallelizable code.*
- This is our new reality!



Problem of Growing!

- We live in the **multi-core/multi-processor era**.
- But we're not prepared for it ...
 - Most of our software is **non-parallelizable**.
 - Most of our software is written for **single-processor**.
 - Most of our software has **shared state**.



Present day approach

- Shared state model.
 - The way we're used to.
 - We have a few variables.
 - We have one or more threads.
 - We have our threads accessing our variables.
 - We have to acquire/release **locks**.
 - ✓ The right locks.
 - ✓ In the right order.
 - ✓ For the right resources.



Difficulties with concurrent execution

Concurrency introduces complexity:

- ▶ concurrent processes may corrupt shared data (Safety)
- ▶ processes may "starve" if not properly coordinated (Liveness)
- ▶ the same program run twice may give different results (Non-determinism)
- ▶ thread construction, context switching and synchronization take time (Run-time overhead)



Concurrency by creating new processes

Remember C library function `fork()` ?

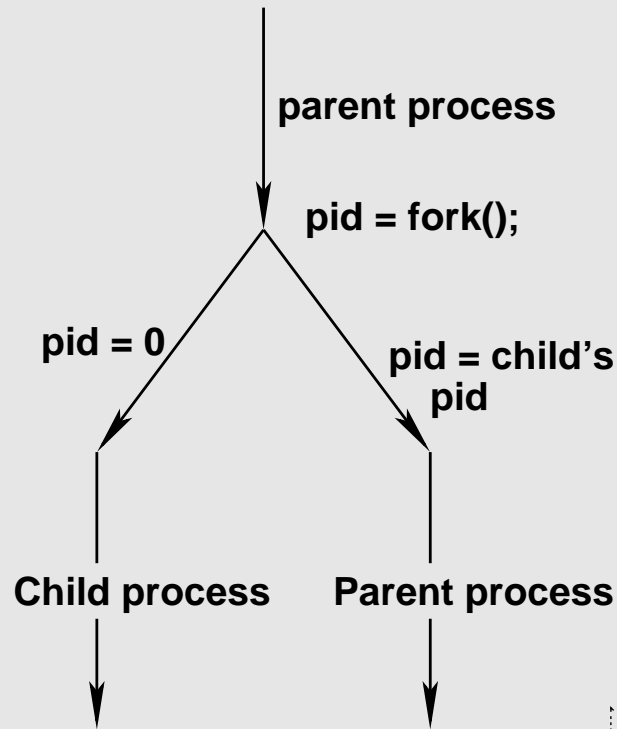


Figure 2



Threads

Consider a file server that has several *"Threads of Control"* or *"Context of Execution"* as Linus calls it.

- ▶ The server will block occasionally for disk I/O.
- ▶ While one thread is sleeping, another could continue its execution.
- ▶ Net result: higher throughput, better performance.




Figure 3

Sharing of Resources

- ▶ Two problems: Race conditions and Dead-locks
- ▶ co-ordination between Threads required:
 - critical sections
 - mutual exclusion



What is the value of x?

Execution-context P1 and P2 run concurrently:

▶ $x \leftarrow 0$

▶ P1: $x \leftarrow x+1$

▶ P2: $x \leftarrow x+2$

▶ $x = ?$

Concurrency and atomicity



Race conditions

Time Counter = 0

Reg ← Counter

Reg ← Reg + 1

Counter ← Reg

Reg ← Counter

Reg ← Reg + 1

Counter ← Reg

Counter = 2

Counter = 0

Reg ← Counter

Reg ← Counter

Reg ← Reg + 1

Reg ← Reg + 1

Counter ← Reg

Counter ← Reg

Counter = 1

What is the value of x?

Execution-context P1 and P2 run concurrently:

$x \leftarrow 0$

P1: $x \leftarrow x+1$

P2: $x \leftarrow x+2$

$x = ?$

$x \leftarrow 0$

P1

$R \leftarrow x$

$R \leftarrow R+1$

$x \leftarrow R$

P2

$R \leftarrow x$

$R \leftarrow R+2$

$x \leftarrow R$

x can be 1, 2 or 3 !!!

Safety

Safety = ensuring **consistency**

- ▶ **Mutual exclusion:** shared resources must be updated atomically
- ▶ **Condition synchronization:** operations may be **delayed** if shared resources are in the wrong state, (e.g., read from empty buffer)



Liveness

Liveness = ensuring **progress**

- ▶ **No Deadlock:** some process can always access a shared resource
- ▶ **No Starvation:** all processes can eventually access shared resources



Expressing Concurrency

- ▶ **Process creation** how do you specify concurrent processes?
- ▶ **Communication** how do processes exchange information?
- ▶ **Synchronization** how do processes maintain consistency?



Concurrent Process Creation

Most concurrent languages offer some variant of the following:

- ▶ Co-routines
- ▶ Fork and Join
- ▶ Cobegin/coend



Co-routines

- ▶ Co-routines are only *pseudo-concurrent* and require *explicit transfers of control*
- ▶ Co-routines can be used to implement most higher-level concurrent mechanisms.

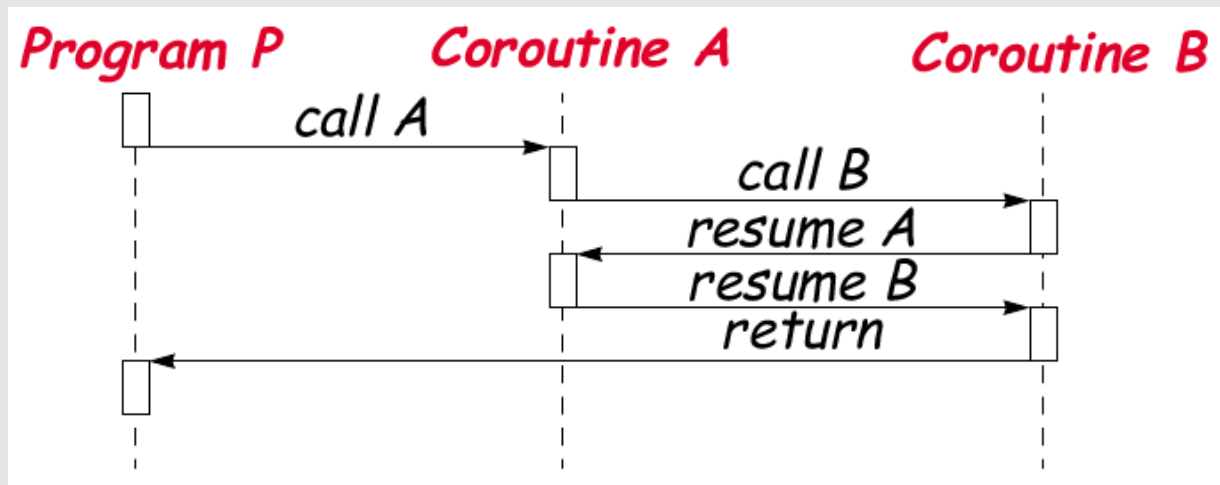


Figure 4

Fork and Join

- ▶ Fork can be used to create any number of processes
- ▶ Join waits for another process to terminate
- ▶ Fork and join are unstructured, so require care and discipline

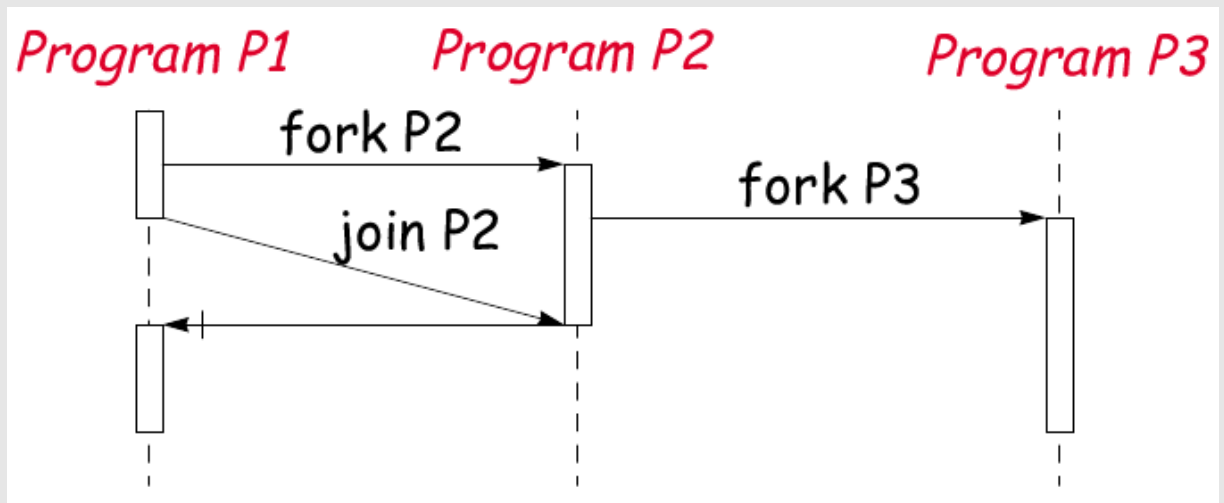


Figure 5



Cobegin/coend

- ▶ Cobegin/coend blocks are better structured
- ▶ `cobegin S1 | S2 | ... | Sn coend`
- ▶ they can only create a fixed number of processes
- ▶ The caller continues when all of the coblocks have terminated





Figure 6




Critical Sections and Mutual Exclusion



Communication and Synchronization


Communication and Synchronization



In approaches based on *shared variables*, processes communicate *indirectly*. *Explicit synchronization mechanisms* are needed.

In *message passing* approaches, *communication and synchronization are combined*.

Communication may be either *synchronous* or *asynchronous*.



Synchronization Techniques

Different approaches are roughly equivalent in expressive power and can be used to implement each other.

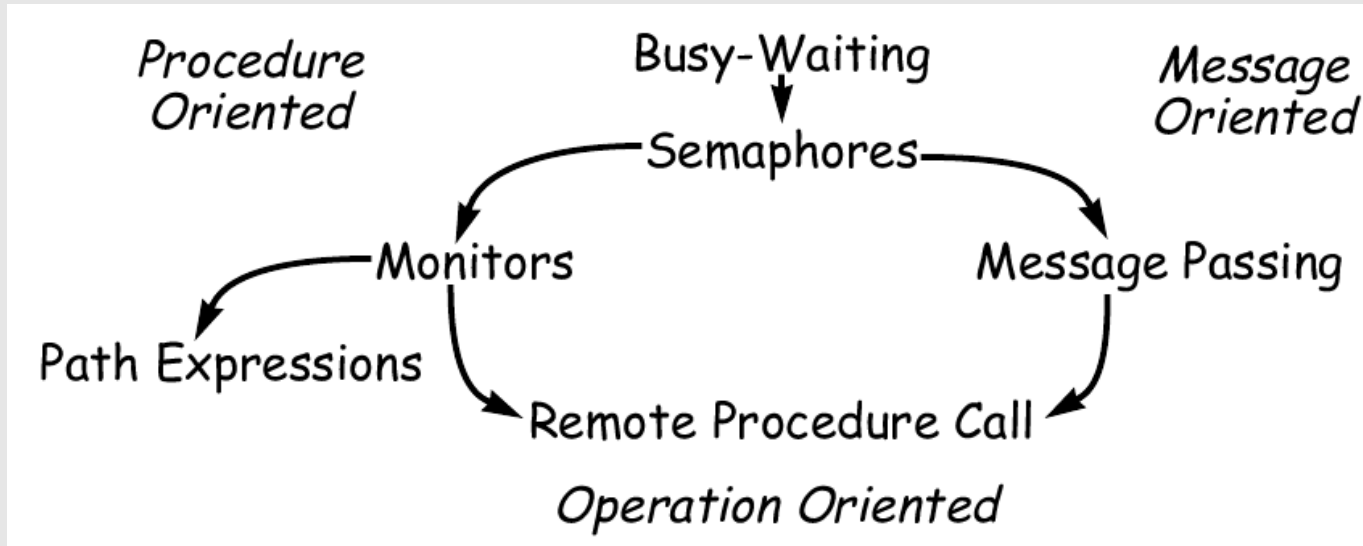


Figure 7

Each approach emphasizes a different style of programming.



Busy-Waiting

- ▶ Busy-waiting is primitive but effective
- ▶ Processes atomically set and test shared variables
- ▶ Condition synchronization is easy to implement
 - to **signal** a condition, a process sets a shared variable
 - to **wait** for a condition, a process repeatedly tests the variable
- ▶ Mutual exclusion is more difficult to realize correctly and efficiently



Semaphores

- ▶ introduced by **E.W. Dijkstra** (1968)
- ▶ a higher-level primitive for process synchronization
- ▶ a non-negative, integer-valued variable s with two operations:
 - **P(s)** - delays (waits) until $s > 0$,
then, atomically executes $s \leftarrow s - 1$
 - **V(s)** - atomically executes $s \leftarrow s + 1$



Programming with semaphores

```
process P1
  loop
    P(mutex)
    Critical Section
    V(mutex)
    Non-critical Sec
  end
end
```

```
process P2
  loop
    P(mutex)
    Critical Section
    V(mutex)
    Non-critical Sec
  end
end
```



Monitors

A monitor encapsulates **resources** and **operations** that manipulate them:

- ▶ operations are invoked like ordinary procedure calls
- ▶ invocations are guaranteed to be **mutually exclusive**
- ▶ condition synchronization is realized using **wait** and **signal** primitives
- ▶ there exist many variations of wait and signal



Other Mechanisms

- ▶ Message Passing: combines communication and synchronization
- ▶ Remote Procedure Calls (RPC)
- ▶ Rendezvous



Concurrency and Parallelism

- ▶ N processes, N processors – Full Parallelism, Concurrency
- ▶ N processes, 1 CPU – no Parallelism, Concurrency
- ▶ N processes, M CPU ($M < N$) – some Parallelism, Concurrency
- ▶ actual single CPU systems – CPU + DMA (a real processor) + Interrupt-driven I/O (illusion of several processors)
- ▶ Virtual Machines – illusion of many CPU's
- ▶ concurrency is independent of parallelism
- ▶ even with a single CPU, even w/o interrupts (consider Time-sharing systems)

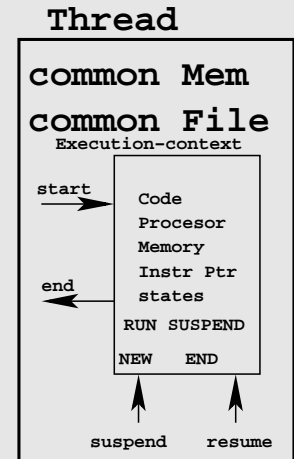
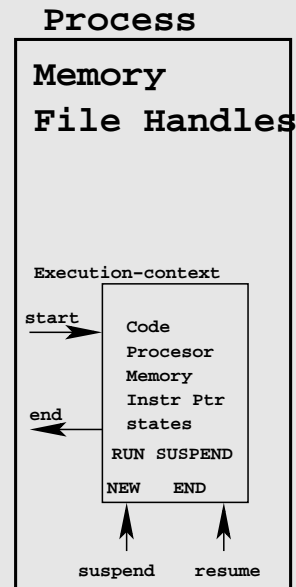
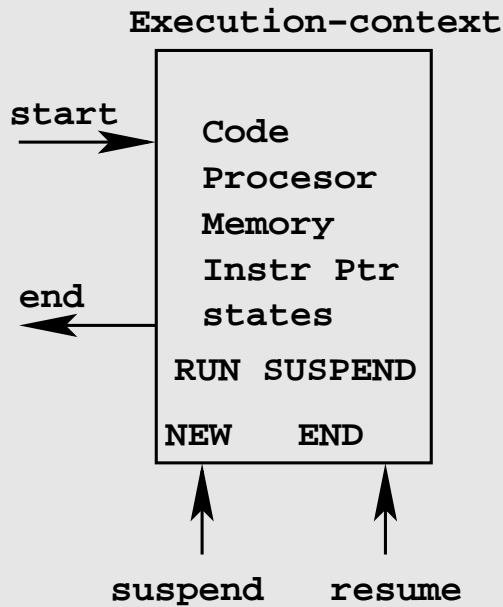


Three basic ideas

- ▶ **communication**: conveying of information from one entity to another
- ▶ **co-ordination**: entities sharing resources are to co-ordinate their activities, for dependable operation
- ▶ **concurrency**: sharing of resurces in the same time frame - CPU, memory, data, code, devices



Execution Contexts



What is a Process?

- ▶ *A process is an instance of a running program*, but more precisely -
- ▶ runs and provides an environment for a program
- ▶ consists of an address space and control point
- ▶ it is basic scheduling entity, only one process runs at a time
- ▶ contends for and owns various system resources like memory
- ▶ requests system services, provided by kernel
- ▶ it has a life-time and Life-cycle
- ▶ is part of a hierarchy - *init* is ancestor of all



Process context

- ▶ User address space: text (code), data, user stack, shared memory, etc.
- ▶ Control information - proc structure, kernel stack, address translation tables
- ▶ Credentials - user and group ID's
- ▶ Environmental variables -
 - a set of strings of the form "var = value"
 - library provides functions to manipulate these
 - inherited from the parent process
 - while invoking a new process, option of new environment
- ▶ Hardware context - IP(PC), SP, registers, PSW, MMU regs, FPU regs



Threads

- ▶ Consider a file server that has several "Threads of Control"
- ▶ The server will block occasionally for disk I/O
- ▶ While one thread is sleeping, another could continue its execution
- ▶ Net result: higher throughput, better performance



Threads - 2

- ▶ sometimes called Light Weight Process (LWP)
- ▶ they are like mini-processes
- ▶ runs strictly sequentially, own PC, stack
- ▶ share CPU in time-share manner
- ▶ only on multiple CPU (multiprocessor) can they really in parallel
- ▶ can create child threads, block on system calls
- ▶ while one thread is blocked, other in the same process can run
- ▶ Creation of a Thread is 10 to 25 times faster than a Processes, e.g. time for 50,000 creations:

Platform	fork()			pthread_create()		
2.4GHz Xeon (2 cpus/node)	real	user	sys	real	user	sys
	54.9	1.5	20.8	1.6	0.7	0.9



Per thread items vs. per process items

Per thread items:

- ▶ Program counter
- ▶ stack, SP
- ▶ Register set (How?)
- ▶ Child threads
- ▶ State

Per process items:

- ▶ Address space, global variables;
- ▶ Open files, Timers, Signals;
- ▶ Child processes





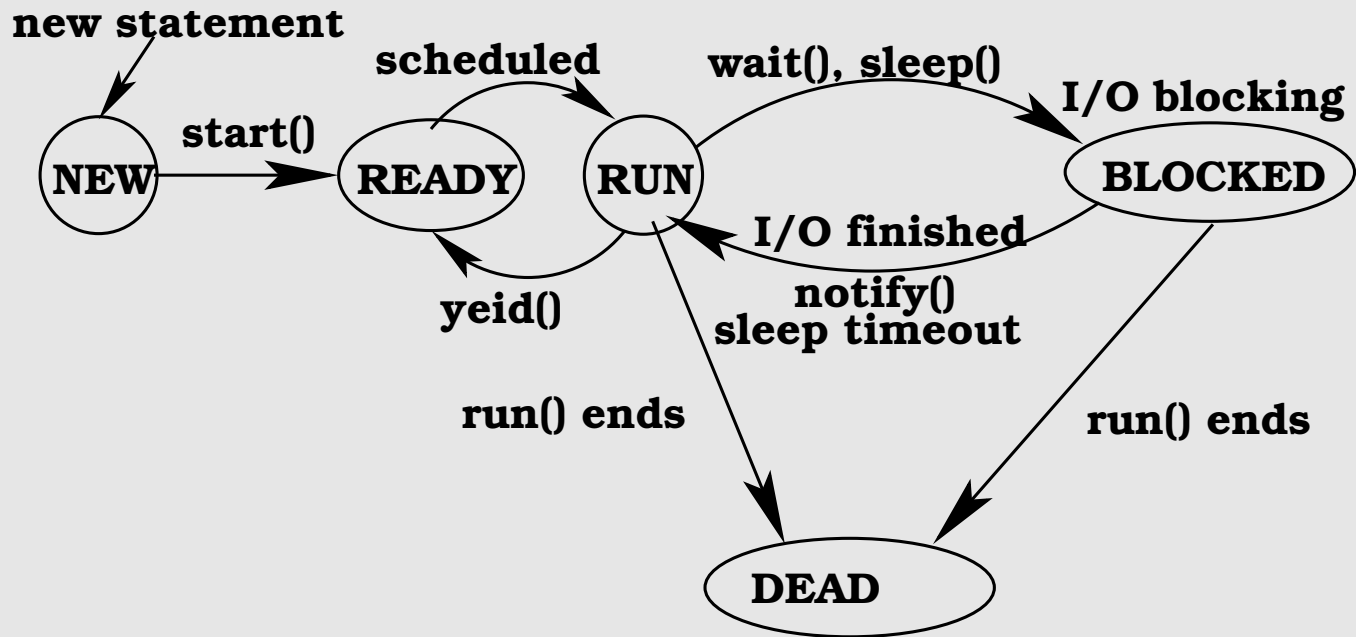
Semaphores



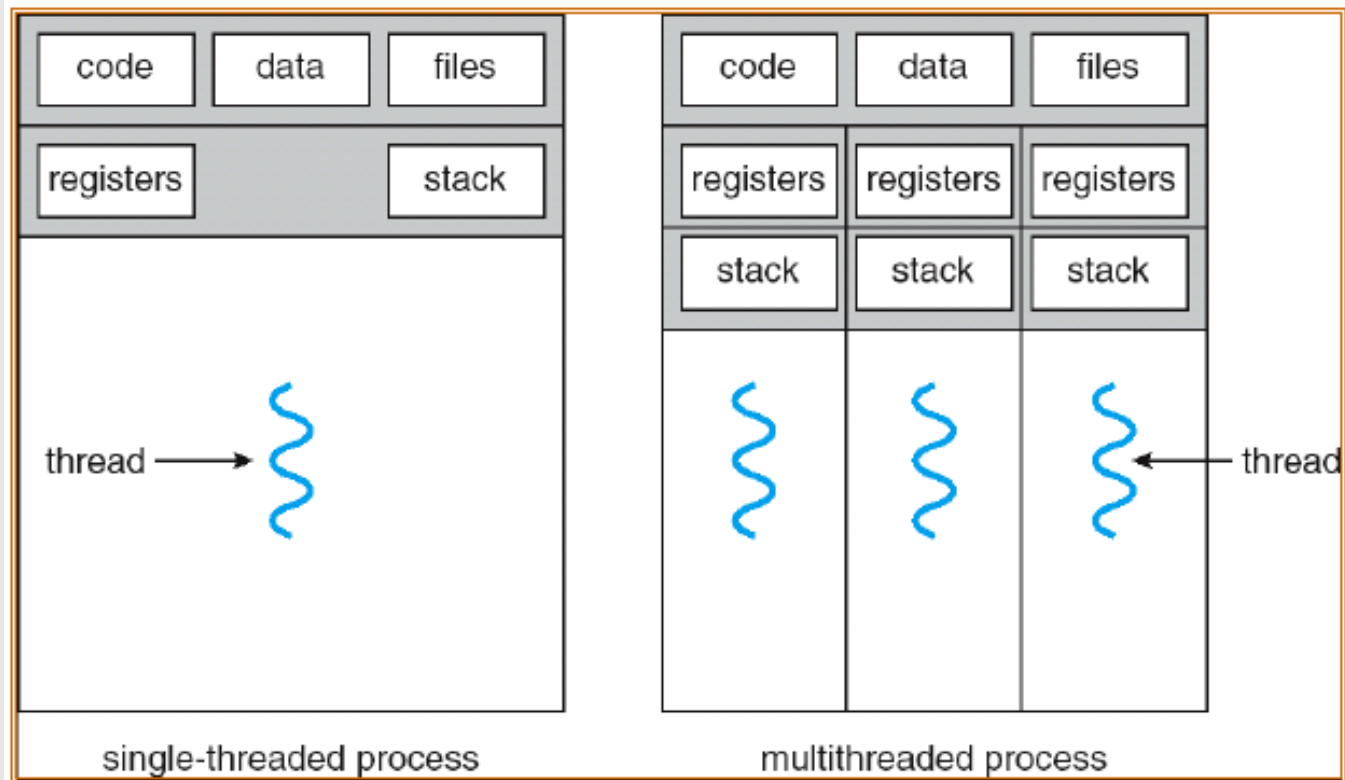
Accounting information



Life-cycle of a Thread



Single and Multithreaded Processes



Single and Multithreaded Processes

- ▶ Threads encapsulate concurrency: **Active** component
- ▶ Address spaces encapsulate protection: **Passive** part
 - Keeps buggy program from trashing the system
- ▶ Why have multiple threads per address space?



On writing concurrent code

- ▶ High-level constructs desirable
- ▶ use low-level constructs like - threads, semaphores, mutex, condition variables
- ▶ `parallel begin`: `initiate`
- ▶ `shared variables`: `shared int i`
- ▶ `Critical Region`: `region i do { ... }`
- ▶ `Conditional Critical Region`: delays a process until components of a shared variable `v` satisfy a condition `B`, `await()`
- ▶ `Semaphore`: - define, init and `Wait()` and `Signal()`
- ▶ we use a pre-processor to map these high-level constructs to functions available in C library